

# College Admission

## DESCRIPTION:

### Background and Objective:

Every year thousands of applications are being submitted by international students for admission in colleges of the USA. It becomes an iterative task for the Education Department to know the total number of applications received and then compare that data with the total number of applications successfully accepted and visas processed. Hence to make the entire process easy, the education department in the US analyze the factors that influence the admission of a student into colleges. The objective of this exercise is to analyse the same.

**Domain:** Education

### Dataset Description:

Attribute	Description
GRE	Graduate Record Exam Scores
GPA	Grade Point Average
Rank	It refers to the prestige of the undergraduate institution. The variable rank takes on the values 1 through 4. Institutions with a rank of 1 have the highest prestige, while those with a rank of 4 have the lowest.
Admit	It is a response variable; admit/don't admit is a binary variable where 1 indicates that student is admitted and 0 indicates that student is not admitted.
SES	SES refers to socioeconomic status: 1 - low, 2 - medium, 3 - high.
Gender_male	Gender_male (0, 1) = 0 -> Female, 1 -> Male
Race	Race – 1, 2, and 3 represent Hispanic, Asian, and African-American

**Analysis Tasks:** Analyze the historical data and determine the key drivers for admission.

## **Source Code:**

```
# ===== Loading Libraries =====  
  
library(readxl) # for reading excel files  
library(caTools) # for splitting dataset  
library(MLmetrics) # for machine learning metrics  
library(caret) # for feature importance  
library(lattice)  
library(tidyverse)  
library(ggpubr) # for creating and customizing 'ggplot2'- based publication ready plots  
library(factoextra)  
library(rpart)  
library(caTools)  
library(randomForest)  
library(kernlab)  
library(readr)  
library(rpart.plot)  
library(naivebayes)  
library(stats)  
  
# ===== Removing Existing R-Objects in Environment =====  
  
rm(list = ls())
```

```
> # loading libraries  
> library(readxl) # for reading excel files  
> library(caTools) # for splitting dataset  
> library(MLmetrics) # for machine learning metrics  
> library(caret) # for feature importance  
> library(lattice)  
> library(tidyverse)  
> library(ggpubr) # for creating and customizing 'ggplot2'- based publ  
ication ready plots  
> library(factoextra)  
> library(rpart)  
> library(caTools)  
> library(randomForest)  
> library(kernlab)  
> library(readr)  
> library(rpart.plot)  
> library(naivebayes)  
> library(stats)  
>  
> # removing objects in environment  
> rm(list = ls())  
> |
```

```
# ===== Loading Dataset =====
```

```
df <- read.csv("College_admission.csv", header = T)
```

```
View(df)
```

```
# ===== EDA =====
```

```
dim(df)
```

```
str(df)
```

```
summary(df)
```

```
~/R prac/ ↗  
> # loading dataset  
> df <- read.csv("College_admission.csv", header = T)  
> View(df)  
>  
>  
> # EDA  
> dim(df)  
[1] 400 7  
> str(df)  
'data.frame': 400 obs. of 7 variables:  
 $ admit : int 0 1 1 1 0 1 1 0 1 0 ...  
 $ gre : int 380 660 800 640 520 760 560 400 540 700 ...  
 $ gpa : num 3.61 3.67 4 3.19 2.93 3 2.98 3.08 3.39 3.92 ...  
 $ ses : int 1 2 2 1 3 2 2 2 1 1 ...  
 $ Gender_Male: int 0 0 0 1 1 1 1 0 1 0 ...  
 $ Race : int 3 2 2 2 2 1 2 2 1 2 ...  
 $ rank : int 3 3 1 4 4 2 1 2 3 2 ...  
> summary(df)  
 admit gre gpa ses  
Min. :0.0000 Min. :220.0 Min. :2.260 Min. :1.000  
1st Qu.:0.0000 1st Qu.:520.0 1st Qu.:3.130 1st Qu.:1.000  
Median :0.0000 Median :580.0 Median :3.395 Median :2.000  
Mean :0.3175 Mean :587.7 Mean :3.390 Mean :1.992  
3rd Qu.:1.0000 3rd Qu.:660.0 3rd Qu.:3.670 3rd Qu.:3.000  
Max. :1.0000 Max. :800.0 Max. :4.000 Max. :3.000  
 Gender_Male Race rank  
Min. :0.000 Min. :1.000 Min. :1.000  
1st Qu.:0.000 1st Qu.:1.000 1st Qu.:2.000  
Median :0.000 Median :2.000 Median :2.000  
Mean :0.475 Mean :1.962 Mean :2.485  
3rd Qu.:1.000 3rd Qu.:3.000 3rd Qu.:3.000  
Max. :1.000 Max. :3.000 Max. :4.000  
> |
```

```
# ===== Question-1 =====
```

```
# checking and handling NA
```

```
sapply(df,function(x) sum(is.na(x))) #no NAs
```

```
~/R prac/
> # checking and handling NA
> sapply(df,function(x) sum(is.na(x))) #no NAs
      admit      gre      gpa      ses Gender_Male
      0        0        0        0        0
      Race      rank
      0        0
> |
```

**From the result we can say that there are no missing values in the data frame.**

```
# ===== Question-2 =====
```

```
# checking for outliers
```

```
boxplot(df)
```

```
boxplot.stats(df$gre)$out
```


```
boxplot.stats(df$gpa)$out
```

```
df <- df[(df$gre >300 & df$gpa != 2.26),]
```

```
boxplot.stats(df$gre)$out
```

```
boxplot.stats(df$gpa)$out
```

```
boxplot(df)
```




**data contains outliers**

**From the figure , we can say that**

```

> boxplot.stats(df$gre)$out
[1] 300 300 220 300
> boxplot.stats(df$gpa)$out
[1] 2.26
> df <- df[(df$gre > 300 & df$gpa != 2.26),]
> boxplot.stats(df$gre)$out
integer(0)
> boxplot.stats(df$gpa)$out
numeric(0)
>

```



**We remove gre data values which are  $\leq 300$  and gpa data value which is 2.26.**

**We only check for gre and gpa features because these are continuous columns else are categorical.**

# creating grade column

```
df$grade <- ifelse(df$gre <= 440, 'Low', ifelse(df$gre > 440 & df$gre <= 580, 'Medium',
ifelse(df$gre > 580, 'High', 'Other'))
```

```
table(df$grade)
```

```
View(df)
```

#plotting barplot for grade and admit side-by-side

```
barplot(table(df$admit, df$grade),
```

```
  beside = T,
```

```
  axisnames = T,
```

```
  xlab = 'Grade-Admit(0/1)',
```

```
ylab = 'Frequency',
main = 'Grade-Admit Chart'
```

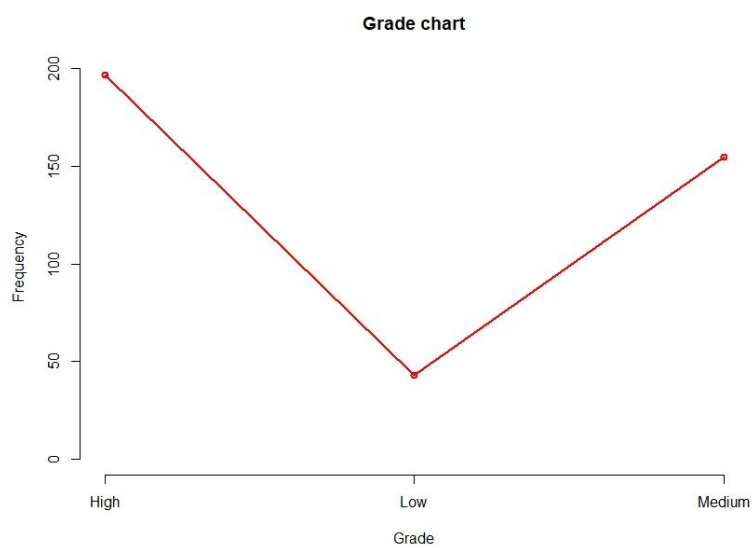
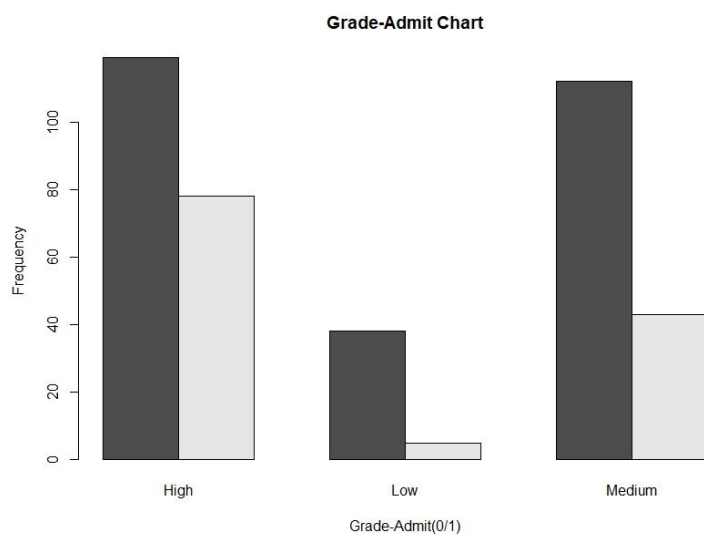
```
)
```

```
#Line Chart for Grade
```

```
plot(table(df$grade),type = "o",col = "red", xlab = "Grade", ylab = "Frequency",
main = "Grade chart")
```

```
~/R prac/ > # creating grade column
> df$grade <- ifelse(df$gre<=440,'Low', ifelse(df$gre>440&df$gre<=580,'Medium', ifelse(df$gre>580,'High','Other')))
> table(df$grade)

  High   Low Medium 
  197    43   155 
> |
```



# ===== Question-3 =====

# structure of data-frame and converting required numeric column to factor and vice-verse

str(df)

df <- df %>% mutate\_at(c(1,5,6), funs(factor(.)))

df\$ses <- factor(df\$ses, ordered = T, levels = c(1:3))

df\$grade <- factor(df\$grade, ordered = T, levels = c('Low','Medium','High'))

df\$rank <- factor(df\$rank, ordered = T, levels = c(4:1))

str(df)

View(df)

```
> # structure of data-frame and converting required numeric column to
  factor and vice-verse
> str(df)
'data.frame':   395 obs. of  8 variables:
 $ admit      : int   0 1 1 1 0 1 1 0 1 0 ...
 $ gre        : int  380 660 800 640 520 760 560 400 540 700 ...
 $ gpa        : num   3.61 3.67 4 3.19 2.93 3 2.98 3.08 3.39 3.92 ...
 $ ses        : int   1 2 2 1 3 2 2 2 1 1 ...
 $ Gender_Male: int   0 0 0 1 1 1 1 0 1 0 ...
 $ Race       : int   3 2 2 2 2 1 2 2 1 2 ...
 $ rank       : int   3 3 1 4 4 2 1 2 3 2 ...
 $ grade      : chr   "Low" "High" "High" "High" ...
> df <- df %>% mutate_at(c(1,5,6), funs(factor(.)))
> df$ses <- factor(df$ses, ordered = T, levels = c(1:3))
> df$grade <- factor(df$grade, ordered = T, levels = c('Low','Medium',
  'High'))
> df$rank <- factor(df$rank, ordered = T, levels = c(4:1))
> str(df)
'data.frame':   395 obs. of  8 variables:
 $ admit      : Factor w/ 2 levels "0","1": 1 2 2 2 1 2 2 1 2 1 ...
 $ gre        : int  380 660 800 640 520 760 560 400 540 700 ...
 $ gpa        : num   3.61 3.67 4 3.19 2.93 3 2.98 3.08 3.39 3.92 ...
 $ ses        : Ord.factor w/ 3 levels "1"<"2"<"3": 1 2 2 1 3 2 2 2 1
  1 ...
 $ Gender_Male: Factor w/ 2 levels "0","1": 1 1 1 2 2 2 2 1 2 1 ...
 $ Race       : Factor w/ 3 levels "1","2","3": 3 2 2 2 2 1 2 2 1 2
  ...
 $ rank       : Ord.factor w/ 4 levels "4"<"3"<"2"<"1": 2 2 4 1 1 3 4
  3 2 3 ...
 $ grade      : Ord.factor w/ 3 levels "Low"<"Medium"<"High": 1 3 3 3 2 3
  2 1 2 3 ...
>
```

**Converting Admit, SES, Race, Rank, Grade features into factors.**

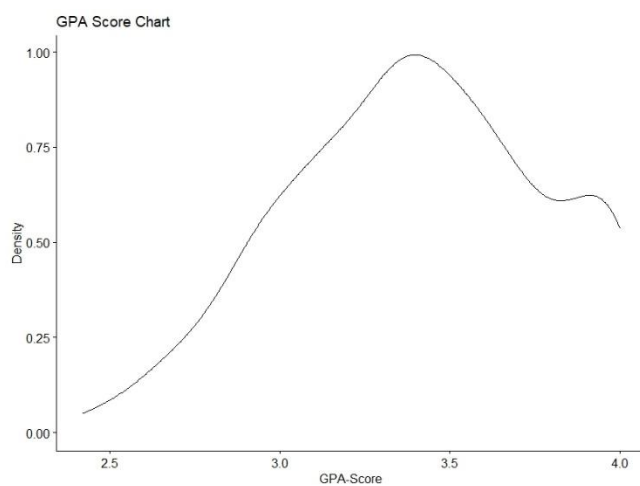
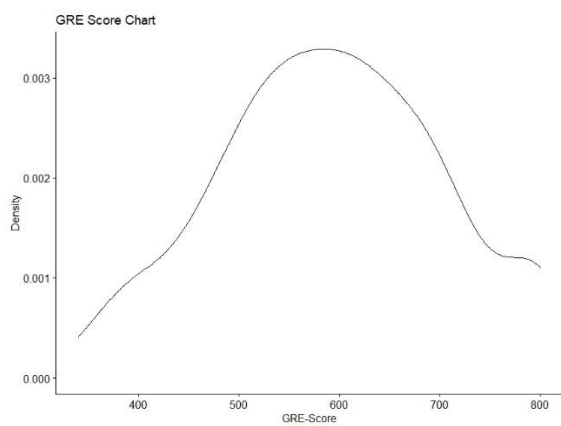
```
# ===== Question-4 =====
```

```
# checking normality
```

```
ggdensity(df$gre,  
          xlab = 'GRE-Score',  
          ylab = 'Density',  
          main = 'GRE Score Chart')
```

```
ggdensity(df$gpa,  
          xlab = 'GPA-Score',  
          ylab = 'Density',  
          main = 'GPA Score Chart')
```

```
sapply(df[2:3], function(x) shapiro.test(x))
```



**From the graphs we can say that data is not normally distributed.**



```

> sapply(df[2:3], function(x) shapiro.test(x))
      gre
statistic 0.9828245
p.value   0.0001223489
method    "Shapiro-Wilk normality test"
data.name "x"
      gpa
statistic 0.9764637
p.value   5.004451e-06
method    "Shapiro-Wilk normality test"
data.name "x"
>

```

**In Shapiro-Wilk Test as p-value is not > 0.05 for both gre and gpa we can say that the data is not normally distributed.**

# ===== Question-5 =====

# Standardization Data

```
df[2:3] <- sapply(df[2:3], function(x) scale(x, center = T, scale = T))
```

```
summary(df)
```

```
View(df)
```

```
mydf <- df #dataframe for decision-tree
```

```

> # Standardization Data
> df[2:3] <- sapply(df[2:3], function(x) scale(x, center = T, sc
T))
> summary(df)
admit      gre      gpa      ses      Gender_Ma
0:269   Min.   :-2.2511   Min.   :-2.604919   1:130   0:209
1:126   1st Qu.: -0.6383   1st Qu.: -0.699963   2:137   1:186
      Median :-0.1007   Median : 0.006071   3:128
      Mean   : 0.0000   Mean   : 0.000000
      3rd Qu.: 0.7057   3rd Qu.: 0.725425
      Max.   : 1.8705   Max.   : 1.604636

Race      rank      grade
1:140   4: 65   Low   : 43
2:128   3:119   Medium:155
3:127   2:150   High  :197
1: 61
>

```

**Scaling data  $Z = \frac{(x - \mu)}{\sigma^2}$  such that mean=0 and standard deviation = 1**

**Hence Data is Normalized using scale().**

# ===== Question-6 =====

# creating dummy variables for factor attributes

```

dummies<- data.frame(sapply(df[,c(1,4:7)],
function(x) data.frame(model.matrix(~x,data =df[,c(1,4:7)]))[,,-1]))

```

```
View(dummies)
```

```
df <- cbind(df[,c(2:3)], dummies)
```

```
# data splitting
```

```
sample <- sample.split(df, SplitRatio = 0.7)
```

```
train <- df[sample,]
```

```
View(train)
```

```
test <- df[!sample,]
```

```
# Feature Importance
```

```
pca <- prcomp(train[, -3])
```

```
summary(pca)
```

```
fviz_eig(pca)
```

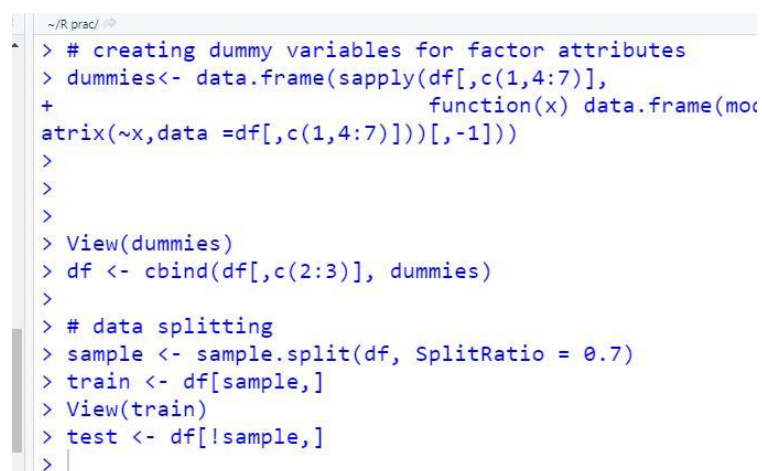
```
fviz_pca_var(pca,
```

```
  col.var = "contrib", # Color by contributions to the PC
```

```
  gradient.cols = c("#00AFBB", "#E7B800", "#FC4E07"),
```

```
  repel = TRUE # Avoid text overlapping
```

```
)
```





```
> # creating dummy variables for factor attributes
> dummies<- data.frame(sapply(df[,c(1,4:7)],
+                           function(x) data.frame(mo:
+ atriX(~x,data =df[,c(1,4:7)])))[-1]))
>
>
>
> View(dummies)
> df <- cbind(df[,c(2:3)], dummies)
>
> # data splitting
> sample <- sample.split(df, SplitRatio = 0.7)
> train <- df[sample,]
> View(train)
> test <- df[!sample,]
> |
```

```

~/R prac/ RPS
> # Feature Importance
> pca <- prcomp(train[, -3])
> summary(pca)
Importance of components:
              PC1      PC2      PC3      PC4      PC5      PC6
Standard deviation  1.1943  0.7439  0.62709  0.58532  0.56777  0.54462
Proportion of Variance 0.3527  0.1368  0.09724  0.08471  0.07971  0.07334
Cumulative Proportion 0.3527  0.4895  0.58676  0.67147  0.75118  0.82452
              PC7      PC8      PC9     PC10
Standard deviation  0.49639  0.44228  0.39627  0.33265
Proportion of Variance 0.06093  0.04837  0.03883  0.02736
Cumulative Proportion 0.88545  0.93381  0.97264  1.00000
> |

```



**From the Figures we can say that GRE and GPA are important features.**

**From the summary of PCA we can say that PC1 to PC6 Contributes around 82% of importance features Cumulatively.**

```
# ===== Question-7 & 8 =====
```

```
# logistic model
```

```
set.seed(123) #for randomness
```

```
#building model
```

```
lmodel <- glm(formula = admit~.,
```

```
          data = train,
```

```
          family = 'binomial')
```

```
summary(lmodel)
```

```
#testing model on test data
```

```
pred <- predict(lmodel,
```

```
          type = 'response',
```

```
          newdata = test[, -3])
```

```
test$actual <- factor(ifelse(test$admit == 1, 'Yes', 'No'))
```

```
test$pred <- factor(ifelse(pred >= 0.5, 'Yes', 'No')) #taking cutoff as 0.5
```

```
View(test)
```

```
#confusion-matrix
```

```
confmatrix <- confmatirx <- confusionMatrix(test$pred,
```

```
          test$actual,
```

```
          positive = 'Yes')
```

```
confmatirx
```

```
#function to get optimal cut-off
```

```
perform_fn <- function(cutoff)
```

```
{
```

```
  pred <- factor(ifelse(pred >= cutoff, "Yes", "No"))
```

```
  conf <- confusionMatrix(pred, test$actual, positive = "Yes")
```

```
  acc <- conf$overall[1]
```

```

sens <- conf$byClass[1]
spec <- conf$byClass[2]
out <- t(as.matrix(c(sens, spec, cutoff, acc)))
colnames(out) <- c("sensitivity", "specificity", 'cutoff', 'accuracy')
return(out)
}

```

```

s = seq(.01,.80,length=100)
s

```

```

OUT = matrix(0,100,4)
OUT
for(i in 1:100)
{
  OUT[i,] = perform_fn(s[i])
}
OUT

```

```

# Let's choose a cutoff value of 0.52070707 for final model
test_cutoff <- factor(ifelse(pred >=0.52070707 , "Yes", "No"))
conf_final <- confusionMatrix(test_cutoff, test$actual, positive = "Yes")
pval <- conf_final$overall[6]
acc <- conf_final$overall[1]
sens <- conf_final$byClass[1]
spec <- conf_final$byClass[2]
pval
acc
sens
spec

```

```
> confmatirx <- confmatirx <- confusionMatrix(test$pred,
+                                             test$actual,
+                                             positive = 'Yes')
> # logistic model
> set.seed(123) #for randomness
> confmatirx
Confusion Matrix and Statistics

          Reference
Prediction No Yes
   No      87  37
   Yes      8  11

      Accuracy : 0.6853
      95% CI : (0.6024, 0.7603)
   No Information Rate : 0.6643
   P-Value [Acc > NIR] : 0.3319

      Kappa : 0.1704

  Mcnemar's Test P-Value : 2.993e-05

   Sensitivity : 0.22917
   Specificity : 0.91579
   Pos Pred Value : 0.57895
   Neg Pred Value : 0.70161
   Prevalence : 0.33566
   Detection Rate : 0.07692
   Detection Prevalence : 0.13287
   Balanced Accuracy : 0.57248

   'Positive' Class : Yes

> # Let's choose a cutoff value of 0.520707 for final model
> test_cutoff <- factor(ifelse(pred >=0.520707 , "Yes", "No"))
> conf_final <- confusionMatrix(test_cutoff, test$actual, positive = "Yes")
>
> pval <- conf_final$overall[6]
> acc <- conf_final$overall[1]
> sens <- conf_final$byClass[1]
> spec <- conf_final$byClass[2]
>
> pval
AccuracyPValue
      0.5424842
> acc
Accuracy
0.7202797
> sens
Sensitivity
      0.4
> spec
Specificity
0.8446602
>
```

**Running the model by randomly taking cut-ff value as 0.5 gives 68% accuracy and Sensitivity=0.22 Specificity=0.91**

**So as to have balance between Sensitivity and Specificity, running model on different cut-off values from 0.01 to 0.8.**

**Choosing the best optimal Cut-off value from those values.**

**Best Optimal Cut-off value for this model based on the dataset is 0.5207. With this accuracy increased from 68.53% to 72.02%.**

```
# ===== Question- 9, 10, 11, 12 =====
```

```
#random forest model
```

```
#train-control for random forest
```

```
control <- trainControl(method="repeatedcv",  
                        number=10,  
                        repeats=3,  
                        savePredictions=TRUE,  
                        classProbs=TRUE,  
                        summaryFunction = twoClassSummary)
```

```
#building model
```

```
rfmodel <- rpart(admit~.,  
               data = train,  
               method = 'class')
```

```
# make predictions on the test set
```

```
predrf <- predict(rfmodel,  
                test,  
                type = 'class')
```

```
test$predrf <- factor(ifelse(predrf == 1,'Yes','No'))
```

```
#confusion-matrix
```

```
rfconmatrix <- confusionMatrix(test$predrf,test$actual,positive = 'Yes')
```

```
rfconmatrix
```

```

+                                     type = class ,
> test$predrf <- factor(ifelse(predrf == 1,'Yes','No'))
> #confusion-matrix
> rfconmatrix <- confusionMatrix(test$predrf,test$actual,positive =
s')
> rfconmatrix
Confusion Matrix and Statistics

```

```

      Reference
Prediction No Yes
      No   87  39
      Yes   6  12

      Accuracy : 0.6875
      95% CI   : (0.605, 0.7621)
      No Information Rate : 0.6458
      P-Value [Acc > NIR] : 0.1691

```

```

      Kappa : 0.2

```

```

McNemar's Test P-Value : 1.84e-06

```

```

      Sensitivity : 0.23529
      Specificity : 0.93548
      Pos Pred Value : 0.66667
      Neg Pred Value : 0.69048
      Prevalence : 0.35417
      Detection Rate : 0.08333
      Detection Prevalence : 0.12500
      Balanced Accuracy : 0.58539

```

```

      'Positive' Class : Yes

```

```

# =====

```

```

#decision tree model

```

```

dttrain <- mydf[sample,]

```

```

dttest <- mydf[!sample,]

```

```

#building model

```

```

dtmodel <- rpart(admit ~ .,
  data = dttrain,
  method = "class",
  control = rpart.control(minsplit = 500,
    minbucket = 250,
    cp = 0.05))

```

```

# make predictions on the test set

```

```

tree.predict <- predict(dtmodel, dttest, type = "class")

```

```

#confusion-matrix

```



```
confusionMatrix(tree.predict, as.factor(dttest$admit), positive = '1')
```

```
> # make predictions on the test set
> tree.predict <- predict(dtmodel, dttest, type = "class")
>
> #confusion-matrix
> confusionMatrix(tree.predict, as.factor(dttest$admit), positive = '1')
Confusion Matrix and Statistics
```

```
      Reference
Prediction 0  1
0    93  51
1     0   0
```

```
      Accuracy : 0.6458
      95% CI : (0.5619, 0.7237)
No Information Rate : 0.6458
P-Value [Acc > NIR] : 0.538
```

```
      Kappa : 0
```

```
McNemar's Test P-Value : 2.534e-12
```

```
      Sensitivity : 0.0000
      Specificity : 1.0000
Pos Pred Value :    NaN
Neg Pred Value : 0.6458
Prevalence : 0.3542
Detection Rate : 0.0000
Detection Prevalence : 0.0000
Balanced Accuracy : 0.5000
```

```
      'Positive' Class : 1
```

```
> |
```

```
# =====
```

```
#naive bayes
```

```
#building model
```

```
nbmodel <- naive_bayes(admit ~ ., data = dttrain, usekernel = T)
```

```
# make predictions on the test set
```

```
nbpred <- predict(nbmodel, dttest)
```

```
#confusion-matrix
```

```
test_conf2 <- confusionMatrix(nbpred, factor(dttest$admit), positive = '1')
```

```
test_conf2
```

```

>
> #confusion-matrix
> test_conf2 <- confusionMatrix(nbpred, factor(dttest$admit),positive =
'1')
> test_conf2
Confusion Matrix and Statistics

          Reference
Prediction 0  1
0      78  35
1      15  16

          Accuracy : 0.6528
          95% CI   : (0.569, 0.7301)
    No Information Rate : 0.6458
    P-Value [Acc > NIR] : 0.46866

          Kappa : 0.1672

    Mcnemar's Test P-Value : 0.00721

          Sensitivity : 0.3137
          Specificity : 0.8387
    Pos Pred Value   : 0.5161
    Neg Pred Value   : 0.6903
    Prevalence       : 0.3542
    Detection Rate   : 0.1111
    Detection Prevalence : 0.2153
    Balanced Accuracy : 0.5762

    'Positive' Class : 1

```

# =====

Parameters/Model	Logistic	Random Forest	Decision Tree	Naïve Baye'
<b>Specificity</b>	0.8446	0.9354	1.0	0.8387
<b>Sensitivity</b>	0.400	0.2352	0	0.3137
<b>P-Value</b>	0.542	0.1691	0.538	0.4686
<b>Accuracy</b>	72.02%	68.75%	64.58%	65.28%

From the above the best model is Logistic Model